

Latent Prediction Is Not Visual Robustness: Diagnosing the Invariance–Resolution Trade-off in JEPA World Models for Control

May 22, 2026

Abstract

Joint-Embedding Predictive Architectures (JEPAs) are commonly *believed* to learn abstract, invariant world representations: by predicting in latent space rather than reconstructing pixels, the encoder is expected to discard visual redundancy and noise. This expectation is a *community heuristic* rather than a published guarantee — to our knowledge no JEPA work formally claims pixel-noise robustness for control. We test the implicit assumption on **LeWorldModel (LeWM)** [22], a published JEPA world model, across four manipulation and navigation tasks (PushT, TwoRoom, Reacher, Cube) and eight levels of train-time pixel-noise augmentation. We find three things:

1. **Visual out-of-distribution (OOD) collapse in JEPA + Cross-Entropy Method (CEM) control.** Without noise-aware training, LeWM collapses under mild pixel noise: PushT control success drops from 86.33% (clean) to 4.67% (Gaussian std = 0.08, near-random); TwoRoom drops from 94.00% to 50.00%.
2. **No globally optimal noise level exists.** Tasks respond very differently to noise augmentation: visually redundant TwoRoom prefers heavy noise (point-best at `std_max` = 0.008), while contact-heavy PushT peaks on *clean* at `std_max` = 0.003 but on *robustness* at `std_max` = 0.006 — clean and robust optima dissociate within a single task.
3. **A five-layer diagnostic protocol explains the compression mechanism.** Instrumenting encoder shift, encoder geometry, predictor sensitivity, latent-noise response, and task resolution traces noise-induced control failure to a representational chain: representation compression (drop in effective rank) → loss of transition-key resolution → loss of controllability (drop in the inverse-dynamics probe R^2). When used as a *cross-checkpoint* predictor, the strongest single diagnostic — which we call the *fragility ratio* (Section 3.3) — tracks *overall checkpoint quality on clean inputs*, not OOD-specific robustness, on the LeWM PushT sweep ($n = 9$). The same partial-correlation null replicates on PLDM PushT ($n = 9$) and on the joint LeWM+PLDM PushT sample ($n = 18$). The toolkit usefully ranks checkpoints by clean control fidelity but does not substitute for actually training with noise when OOD robustness is the goal.

This paper does not propose a new training algorithm. Its contribution is (i) a systematic empirical study of JEPA + CEM visual OOD failure, (ii) a reproducible diagnostic toolkit, and (iii) an explicit delineation of what cross-checkpoint diagnostics can and cannot predict.

Keywords: world models; JEPA; visual robustness; representation diagnostics; invariance–resolution trade-off.

1 Introduction

1.1 The implicit invariance heuristic and where it breaks

Since Yann LeCun proposed JEPA [20], the paradigm has been advanced as a direction for self-supervised learning. Unlike generative models (variational autoencoders, diffusion), JEPA does not reconstruct pixels; it predicts future *representations* in latent space. The motivating intuition — that predicting “what is invariant” rather than “what pixels look like” yields abstract representations that discard visual redundancy and noise — is now part of the field’s informal vocabulary in talks, blog posts, and survey articles [2, 4].

This is a heuristic rather than a published guarantee. To our knowledge no JEPA paper has *formally claimed* visual-OOD robustness for control-relevant downstream tasks. I-JEPA [2] and V-JEPA [4, 3] established strong visual representations on ImageNet and video via masked prediction; LeWorldModel (LeWM) [22] extended the framework to end-to-end stable world-model training across four robotic control tasks. Existing robustness studies probe JEPA only on image classification [24], synthetic 1D distractors [15], or medical ultrasound [25]; none on *JEPA-based control* under realistic pixel noise.

This leaves a basic operational question open: **if the input image is degraded by sensor noise, lighting change, or camera jitter, does a JEPA + CEM world model still plan and act reliably?**

The empirical answer is no. On PushT (2D pushing), the untrained-with-noise LeWM achieves 86.33% on clean images but falls to 4.67% under Gaussian pixel noise of $\text{std} = 0.08$ — near-random. TwoRoom (2D navigation) drops from 94.00% to 50.00%. Latent prediction alone does not, in this regime, confer the visual robustness the community heuristic would predict.

1.2 The core tension: no globally optimal noise level

A natural remedy is input-side noise augmentation during training, a technique long-validated in supervised and contrastive learning [6, 7]. But a deeper question arises: **does there exist a single, universal noise level that is optimal across tasks?**

A systematic sweep across four tasks at eight levels of $\text{std_max} \in \{0.001, \dots, 0.008\}$ answers no:

- **TwoRoom** (visually redundant navigation): clean success rises monotonically with noise, peaking at $\text{std_max} = 0.008$ (98.33% / 98.67%).
- **PushT** (contact-heavy manipulation): clean peaks at $\text{std_max} = 0.003$ (89.67%), but robustness at px+goal 0.08 peaks at $\text{std_max} = 0.006$ (87.00%) — clean and robust optima dissociate.
- **Reacher** (continuous reaching): lies on a 0.002–0.006 plateau; clean point-best is at $\text{std_max} = 0.006$ (86.00%), while px+goal 0.08 point-best is at $\text{std_max} = 0.002$ (85.67%). Very low noise (0.001) is statistically indistinguishable from base.
- **Cube** (structured 3D manipulation): noise sweep is the weakest of the four, with no monotone trend on clean. We read this as the trade-off’s mildest manifestation rather than an absence: a structured action sequence and predictable visual-action coupling supply enough robustness that input-side augmentation buys comparatively little.

This exposes a fundamental tension: **global noise augmentation cannot distinguish “background visual redundancy that should be made invariant” from “control-relevant features that should retain resolution”.**

1.3 Contributions

C1. A systematic quantification of visual-OOD fragility of JEPA + CEM control across four representative tasks. We sweep 4 tasks $\times$ 8 noise levels, and evaluate all 36 checkpoints under a unified 3-evaluation-seed $\times$ 100-trajectory protocol, for 300 trajectories per condition. The same sweep replicates on a second method family (PLDM), confirming that the per-task signatures of the failure mode — visually redundant TwoRoom rises monotonically with noise, contact-heavy PushT has a steep clean-trained cliff with large noise-training recovery, structured 3D Cube responds weakly — are not LeWM-specific (Section F).

C2. The “invariance–resolution trade-off” concept and a five-layer diagnostic protocol that operationalises it. The protocol covers encoder shift, encoder geometry, predictor sensitivity, latent-noise response, and task resolution, and bundles 17 concrete metrics (Section 3.3). We validate it by Spearman correlations on the 4-task $\times$ $n = 9$ LeWM sweep under unified 3-seed $\times$ 100 evaluation, with partial correlations conditioned on training noise to separate residual checkpoint-quality signal from the monotone sweep trend induced by `std_max`.

C3. A mechanistic account of why global noise augmentation has limited returns. On TwoRoom (visually redundant) the gains coincide with desirable representation compression — effective rank drops by roughly 30% ($47.60 \rightarrow 33.59$) and the more compact latent is easier to plan in. On PushT (contact-heavy), the base representation already has higher rank and higher controllability (effective rank 76.42; inverse-dynamics probe $R^2 \approx 0.77$); even the light representative noise-trained checkpoint compresses rank by about 44% ($76.42 \rightarrow 42.85$) with concomitant downward movement in transition-resolution metrics and the controllability probe (both drop by a few points; see Table 4), so further compression cannot be assumed safe for contact-heavy tasks.

C4. An explicit delineation of cross-checkpoint diagnostics. The strongest cross-checkpoint diagnostic (the fragility ratio) carries a residual checkpoint-quality signal beyond the sweep-level `std_max` effect (partial Spearman $\rho = -0.59$ on clean and -0.41 on px+goal 0.08 after conditioning on `std_max`, on the $n = 9$ LeWM PushT sweep with unified 3-seed $\times$ 100 eval). However, the apparent $\rho(\text{metric}, \text{OOD drop}) = -0.77$ is mediated by `std_max`: after partialling out `std_max`, that correlation collapses to $+0.06$. The same null result reappears on PLDM PushT ($\rho = -0.78$ unconditional, -0.14 partial on `std_max`, $n = 9$) and in the joint LeWM+PLDM $n = 18$ sample ($+0.11$ partial on `std_max` and method), so the negative finding is method-robust. The metric is a model-selection tool after controlling for the sweep-level `std_max` trend, not an OOD oracle (Sections 4.5, 5.3 and F).

1.4 Organisation

Section 2 reviews related work; Section 3 defines the LeWM background, noise protocol, and diagnostic framework; Section 4 reports the experimental findings; Section 5 discusses mechanism and implications; Section 6 concludes.

2 Related Work

2.1 JEPA and latent world models

JEPA [20] predicts in latent space rather than reconstructing pixels. I-JEPA [2] predicts target representations from a masked context; V-JEPA [4, 3] extends the framework to video understanding and video-driven world modelling; LeWM [22] is the first end-to-end stable JEPA world model, using **SIGReg** (Sketch Isotropic Gaussian Regularizer) — random projections plus Epps–Pulley empirical-characteristic-function matching [8] — to prevent representation collapse without batch normalisation. LeWM is our baseline. The original LeWM paper reports a Violation-of-Expectation experiment showing the model is sensitive to physical perturbations (object teleportation) but not to visual perturbations (colour change). VoE however measures prediction error (surprise), not control success rate, and colour change differs from pixel-level Gaussian noise. We give the first picture, to our knowledge, of LeWM’s control success rate under pixel-noise corruption.

For a second method family we evaluate PLDM [27, 28] as implemented in `stable-worldmodel` [21] on the full sweep grid (Section F). The PLDM sweep covers the same four tasks and the same eight values of `std_max` as our LeWM sweep, under the same 3-seed $\times$ 100-trajectory evaluation protocol. We use PLDM to test which findings replicate across method families and which are LeWM-specific; the trade-off concept and the diagnostic toolkit themselves are introduced and analysed on LeWM, and the LeWM canonical tables in the main text are kept method-pure for clarity.

2.2 Robustness studies of JEPA

N-JEPA [24] introduces diffusion-noise augmentation into I-JEPA, improving linear-probing robustness on ImageNet. VJEPA [15] tests a Noisy-TV distractor on synthetic 1D signals and reports JEPA retains $R^2 > 0.84$ under high noise. US-JEPA [25] tests Gaussian blur, contrast reduction, and speckle noise on medical ultrasound. The contemporaneous Bisim-JEPA [32] attacks the related problem of slow-feature distractors in JEPA-based planning by augmenting the predictive objective with a bisimulation encoder that maps states with similar dynamics to nearby latents. Their intervention (bisimulation regularizer for control-relevant invariance) is complementary to ours: we instrument an unmodified LeWM under pixel noise and report what its latents *do and do not* retain, while they propose an inductive bias that targets a different corruption family (slow-feature distractors). None of these studies the pixel-noise robustness of a JEPA *world model* on *robotic control* with a cross-checkpoint diagnostic–evaluation correlation analysis. Furthermore, VJEPA’s optimistic conclusion ($R^2 > 0.84$ in 1D) contrasts with our control-time observation (success rate $\rightarrow 4.67\%$), suggesting that the “natural robustness” of JEPA may be an artefact of evaluation modality.

2.3 World models and input augmentation

Input-side image augmentation is an established route to visual robustness in pixel-based reinforcement learning: DrQ [18] and DrQ-v2 [34] regularise a model-free critic with random-shift augmentation; SODA / DMC-GB [14] formalises a visual-generalisation benchmark on DeepMind Control with random colours and video backgrounds. Both lines establish that augmentation pressure during training is a strong baseline for closing the OOD gap — but they study model-free policies with explicit reward signals, not JEPA-style latent-prediction world models with planning-time CEM control. In the model-based RL world-model literature, DreamerV3 [12] and TD-MPC2 [13] rely on learned visual encoders and latent dynamics, which may provide some implicit tolerance to benign visual variation but do not by themselves settle sensor-noise robustness. ViGMO [23] tests Gaussian

noise and blur on DMC tasks, finds “sensor noise is a fundamentally different distribution shift”, and proposes a latent-consistency loss. ViGMO addresses model-based RL, not JEPA architectures; its conclusion is directionally aligned with ours, but our contribution adds mechanistic decomposition via the five-layer diagnostic, not just performance reporting.

2.4 The invariance–resolution tension

The idea that representation learning balances two competing properties is well established in the self-supervised literature. Wang and Isola [33] decompose the contrastive loss into an *alignment* term (positive pairs close together) and a *uniformity* term (features spread on the hypersphere), and show that downstream linear-probe accuracy tracks both in concert. Tamkin et al. [30] argue, in the contrastive-learning context, that label-destroying augmentations can be useful — augmentations act as feature dropout rather than pure invariance inducers. Zhang and Ma [35] note that augmentation modules can impose invariances that must be matched to the task. Concurrent seq-JEPA [11] resolves a related *invariance–equivariance* tension in JEPA by introducing architectural disentanglement of two heads. Our contribution is conceptually adjacent but operationally distinct: we study a *compression-versus-control-resolution* trade-off observed empirically in latent predictive world models, where the downstream task is planning rather than discrimination, and we operationalise it through a diagnostic protocol rather than a new architecture or loss.

2.5 Representation diagnostics and collapse analysis

Self-supervised learning broadly uses effective rank [26], condition number, and participation ratio to diagnose dimensional collapse [16]. Next-Latent Prediction [31] uses effective latent rank to assess world-model compactness. VICReg [5] provides an anti-collapse regularizer distinct from SIGReg. Centered Kernel Alignment (CKA) [17] compares representations across networks or training conditions; we use it to compare clean-input and noisy-input latents within a checkpoint. Linear probes for representation analysis follow Alain and Bengio [1]. Most relevantly, RankMe [10] showed that a label-free effective-rank diagnostic correlates with downstream linear-probe transfer across hyperparameter sweeps in self-supervised learning. Our cross-checkpoint correlation analysis (Section 4.5) asks the analogous question for control: *does a label-free predictor-sensitivity diagnostic correlate with downstream task success after the sweep-level training trend is removed?* Individual metrics are not new; our contribution is to combine them into a coherent five-layer protocol with per-token noise sensitivity, cross-checkpoint correlation, and a partial-correlation validation scheme conditioning on training noise — and to delineate where that programme succeeds and where it does not.

3 Background and Diagnostic Framework

3.1 LeWorldModel baseline

LeWM [22] is an end-to-end JEPA world model trained with two terms only:

$$\mathcal{L}_{\text{LeWM}} = \mathcal{L}_{\text{pred}} + \lambda \cdot \mathcal{L}_{\text{SIGReg}}, \quad (1)$$

where $\mathcal{L}_{\text{pred}}$ is the latent-space mean-squared error (MSE) between predicted and target representations. SIGReg projects each latent onto M unit-norm random directions, computes the Epps–Pulley empirical-characteristic-function distance [8] between each projection and $\mathcal{N}(0, 1)$, and aggregates with Gauss-window weights, preventing collapse without explicit BatchNorm. (The Cramér–Wold

theorem motivates the construction: equality of high-dimensional distributions reduces to equality of all one-dimensional projections of their characteristic functions.) Inference uses the Cross-Entropy Method (CEM) [19] for latent-space model predictive control (MPC) [9].

3.2 Input-side noise augmentation

We add per-frame Gaussian noise to the LeWM input pipeline via `utils.py::AddNormalizedGaussianNoise` (Section A). Each frame is independently noised: $\text{Bernoulli}(p)$ decides whether the frame is corrupted, and if so the standard deviation is drawn from $\text{Uniform}(0, \text{std_max})$. We fix $p = 1.0$ and sweep $\text{std_max} \in \{0.001, \dots, 0.008\}$ (eight levels). Evaluation uses two noised conditions: pixels+goal 0.05 and pixels+goal 0.08 (Gaussian noise of the indicated std applied to both pixels and the goal image).

3.3 Five-layer diagnostic framework

The framework decomposes the JEPA + CEM forward pass — pixels $\rightarrow$ encoder $f \rightarrow$ latent $z \rightarrow$ predictor $g \rightarrow$ cost surface $\rightarrow$ planned actions — into five sequential stages and instruments each transition with one or more metrics. Layers 1–2 characterise the encoder output (how the latent shifts under noise; how the latent space is globally organised); Layer 3 measures how the predictor amplifies that shift; Layer 4 isolates the predictor and cost surface by injecting noise directly into the latent; Layer 5 quantifies the planning-relevant information that the latent still carries. Each layer probes one transition, so an observed failure can be localised to a single stage. Most individual metrics already have a literature home (cited inline below); the contribution is the composite protocol — and a handful of metrics introduced for this paper (the *fragility ratio*, `transition_resolution_ratio`, `latent_robust_radius_z`) that explicitly normalise predictor and cost-surface sensitivity by the local nearest-neighbour scale of the clean encoder. We use nearest-neighbour distance as a local latent scale primitive in the spirit of kNN-based OOD detection [29], but the scale-normalised ratios below are introduced for this diagnostic protocol.

Minimal notation. Let $z_i = f(x_i)$ and $\tilde{z}_i = f(\tilde{x}_i)$ be clean and noised latents for the same token, and let $d_{\cos}(a, b) = 1 - \frac{a^\top b}{\|a\| \|b\|}$. Define the clean local scale $d_i^{\text{NN}} = \min_{j \neq i} d_{\cos}(z_i, z_j)$. The main introduced ratios are:

$$r_i^{\text{enc}} = \frac{d_{\cos}(z_i, \tilde{z}_i)}{d_i^{\text{NN}} + \epsilon}, \quad r_i^{\text{pred}} = \frac{d_{\cos}(g(z)_i, g(\tilde{z})_i)}{d_i^{\text{NN}} + \epsilon}, \quad (2)$$

$$\text{transition_resolution_ratio}_d = \frac{\text{median}_t d(z_t, z_{t+1})}{\text{median}_{(t, t') \in \mathcal{F}} d(z_t, z_{t'}) + \epsilon}, \quad (3)$$

where r_i^{enc} corresponds to `noise_to_nn_cos_ratio`, r_i^{pred} to the fragility ratio, d is either cosine or L2 distance for `transition_resolution_ratio`, and $\mathcal{F}$ denotes random far pairs from different temporal locations. Layer 1 otherwise uses standard Euclidean/cosine displacements; no method-specific external citation is needed for the raw angle or L2 quantities.

Layer 1 — Encoder shift. The magnitude and direction of latent displacement under input noise. Metrics: `noise_angle_deg`, `noise_l2`, `noise_to_nn_cos_ratio` (encoder shift normalised by the clean nearest-neighbour (NN) cosine distance — introduced here), `noise_angle_slope` (the per-unit-std angular gain).

Layer 2 — Encoder geometry. The global structure of the encoded latent space. Metrics: `clean_nn_cos_dist` (local neighbourhood scale), `clean_effective_rank` [26] (intrinsic dimensionality), and `cka_linear_at_max_std` (CKA [17] between clean-input and noisy-input latents at the diagnostic’s max injection std).

Layer 3 — Predictor sensitivity. How the predictor amplifies the encoder shift. Metrics: the *fragility ratio* (single-step predictor target shift normalised by clean NN distance, stored as `predictor_target_to_nn_cos_ratio_at_max_std` — introduced here; the strongest cross-checkpoint diagnostic we identify), and `predictor_rollout_drift_T(T)` (multi-step rollout L2 drift at horizon T).

Layer 4 — Latent-noise response. Noise injected directly into the latent z (bypassing the encoder) isolates predictor and cost contributions. Metrics: `latent_cost_surface_slope_z` (slope of the planner’s cost surface vs latent perturbation magnitude) and `latent_robust_radius_z` (the largest perturbation that keeps the cost ranking stable — introduced here).

Layer 5 — Task resolution. The control-relevant information retained in the latent. Metrics: `transition_resolution_ratio_l2` and `transition_resolution_ratio_cos` (the ratio of inter-transition to intra-transition distance under L2 / cosine geometries — introduced here as a planning-relevant analogue of class separability), and `id_probe_r2`, the inverse-dynamics linear-probe R^2 (a controllability proxy in the spirit of linear-probe analysis [1]).

3.4 Cross-checkpoint validation protocol

To rule out spurious correlations we adopt two complementary analyses on the same LeWM PushT sweep:

LeWM $n = 9$ sweep with partial correlation. All 9 LeWM checkpoints per task (`base + std_max` $\in \{0.001, \dots, 0.008\}$). Compute Spearman ρ and *partial* Spearman conditioned on `std_max`. The partialling step matters: many diagnostics correlate with control performance only because both quantities co-vary with the training noise level along the sweep. The relevant question is whether a diagnostic retains a residual checkpoint-quality signal after removing the monotone sweep trend associated with `std_max`.

We report both the raw Spearman and the partial-on-`std_max` quantities. The raw ρ tells you “which checkpoint over the entire sweep behaves better”; the partial ρ tells you whether a diagnostic retains a residual checkpoint-quality signal after removing the monotone trend associated with `std_max`. The two questions are distinct, and we will see in Section 4.5 that they have markedly different answers for the same metric. A larger cross-architecture protocol (varying the latent geometry of the world model itself) is left to follow-up work with external baselines.

4 Experiments

4.1 Setup

Tasks. PushT (2D pushing), TwoRoom (2D navigation), Reacher (2D arm), Cube (3D manipulation).

Baselines. LeWM-base (no noise) and LeWM+noise (8-level sweep).

Checkpoints. The 36 checkpoints in this paper correspond to one trained model per (task, std_max) configuration.

Evaluation seeds. Each checkpoint is evaluated with 3 evaluation seeds (42 / 43 / 44), with 100 trajectories per seed.

Evaluation protocol. All success rates in this paper — clean and noised, across all 36 checkpoints ($4 \text{ tasks} \times \{\text{base, std_max } 0.001\text{--}0.008\}$) — are computed under a single protocol: $n = 3$ seeds (42/43/44), num_eval = 100 trajectories per seed (300 trajectories per condition per checkpoint). Every cell of Tables 1, 2, 3 is mean $\pm$ across-seed population std over $n = 3$ and is sourced from the released aggregate, assets/paper1_data/canonical_evals_20260517.json; raw per-seed files live alongside each checkpoint as <ckpt>/eval_results/<cond>_seed{42,43,44}_metrics.txt. The diagnostic source-of-truth for Figure 5, Table 5, Table 6, and Table 7 is assets/paper1_data/canonical_diagnostics_20260517.json.

Hardware. Single NVIDIA H800 (80 GB).

Figures. Six main figures (Figures 1 to 6) are produced by tools/paper1_figs.py.

4.2 JEPA control fragility under visual OOD

Table 1 reports LeWM-base success rates under clean and noised evaluation (mean $\pm$ std across 3 seeds $\times$ 100 evaluations).

Table 1: LeWM-base under visual OOD (mean $\pm$ std; 3 seeds $\times$ 100 evaluations).

Task	clean	px+goal 0.05	px+goal 0.08	clean $\rightarrow$ 0.08 drop
TwoRoom	94.00 $\pm$ 3.56	61.33 $\pm$ 5.31	50.00 $\pm$ 1.41	−44.00
PushT	86.33 $\pm$ 2.36	12.00 $\pm$ 4.55	4.67 $\pm$ 2.05	−81.67
Reacher	58.67 $\pm$ 1.25	27.00 $\pm$ 5.10	15.00 $\pm$ 2.16	−43.67
Cube	66.67 $\pm$ 2.62	53.33 $\pm$ 3.30	46.33 $\pm$ 3.68	−20.33




Figure 1: Visual OOD cliff in LeWM and recovery by noise training (success rate, 4 tasks). Blue: LeWM-base clean. Red: LeWM-base under px+g 0.08 noise. Green: the px+g 0.08 point-best LeWM+noise configuration. Annotations show base drop and recovery in points; σ^* marks each task’s px+g 0.08 point-best `std_max`.

LeWM-base is strong on clean images (especially TwoRoom and PushT), but visual `std = 0.05` applied to pixels and goal jointly already produces large drops on all tasks: PushT loses 74+ pts (down to near-random 4.67% at `std = 0.08`), TwoRoom 30+ pts, Reacher 30+ pts, Cube ~ 13 pts (at `std = 0.05`). **This is not a marginal phenomenon:** a JEPA + CEM world model without noise-aware training has essentially no resistance to visual corruption. The drop pattern across tasks is informative: Cube degrades least (-20.33 pt at `std = 0.08`) — structured manipulation has some natural robustness — while PushT degrades most catastrophically (-81.67 pt), confirming that contact-heavy continuous control is most sensitive to visual precision.

The same direction of failure replicates on PLDM: a clean-trained PLDM checkpoint on PushT drops from $75.33\% \pm 3.68$ clean to $10.00\% \pm 2.16$ at px+goal 0.08 (-65.33 pt, Table 13), and the per-task signatures of the noise-training sweep — TwoRoom’s monotone rise to high `std_max`, PushT’s large `std_max` ≈ 0.003 recovery, Cube’s weak response — carry over to the full PLDM sweep that we report in Section F.

4.3 Noise augmentation closes the gap — at the cost of task-specific tuning

Tables 2 and 3 report the complete 8-level sweep across tasks. All cells are success rate $\pm$ across-seed `std` (in pts), aggregated under the unified $n = 3$ seeds (42/43/44) $\times$ 100 trajectories protocol — 300 trajectories per cell.


Table 2: LeWM+noise sweep, **clean** success rate (%).

std_max	TwoRoom	PushT	Reacher	Cube
0 (base)	94.00 $\pm$ 3.56	86.33 $\pm$ 2.36	58.67 $\pm$ 1.25	66.67 $\pm$ 2.62
0.001	93.67 $\pm$ 3.30	88.00 $\pm$ 3.74	61.67 $\pm$ 2.49	69.33 $\pm$ 0.47
0.002	95.00 $\pm$ 2.83	88.33 $\pm$ 2.62	85.67 $\pm$ 2.49	60.00 $\pm$ 1.63
0.003	96.33 $\pm$ 3.30	89.67 $\pm$ 1.70 [†]	78.67 $\pm$ 1.25	65.00 $\pm$ 1.63
0.004	96.33 $\pm$ 2.05	89.33 $\pm$ 2.05	84.00 $\pm$ 2.94	69.00 $\pm$ 3.74
0.005	96.00 $\pm$ 2.83	80.67 $\pm$ 4.78	70.00 $\pm$ 2.16	59.33 $\pm$ 0.94
0.006	96.67 $\pm$ 2.05	89.33 $\pm$ 2.05	86.00 $\pm$ 2.94 [†]	66.67 $\pm$ 2.05
0.007	96.00 $\pm$ 1.63	85.67 $\pm$ 3.09	83.67 $\pm$ 3.30	67.67 $\pm$ 0.94
0.008	98.33 $\pm$ 0.47 [†]	88.33 $\pm$ 2.87	84.00 $\pm$ 0.82	62.33 $\pm$ 1.25

Table 3: LeWM+noise sweep, **px+goal 0.08** success rate (%).

std_max	TwoRoom	PushT	Reacher	Cube
0 (base)	50.00 $\pm$ 1.41	4.67 $\pm$ 2.05	15.00 $\pm$ 2.16	46.33 $\pm$ 3.68
0.001	87.67 $\pm$ 1.89	43.33 $\pm$ 3.09	46.00 $\pm$ 1.63	51.33 $\pm$ 5.79
0.002	93.33 $\pm$ 0.94	71.33 $\pm$ 3.68	85.67 $\pm$ 1.70 [†]	60.67 $\pm$ 0.47
0.003	94.67 $\pm$ 2.87	83.00 $\pm$ 3.74	73.67 $\pm$ 0.47	67.33 $\pm$ 1.89
0.004	95.00 $\pm$ 2.45	81.33 $\pm$ 2.87	80.00 $\pm$ 1.41	67.00 $\pm$ 3.56
0.005	95.67 $\pm$ 2.36	75.00 $\pm$ 6.48	68.00 $\pm$ 3.56	59.67 $\pm$ 2.05
0.006	96.67 $\pm$ 2.49	87.00 $\pm$ 3.74 [‡]	84.67 $\pm$ 4.03	65.00 $\pm$ 2.94
0.007	96.33 $\pm$ 2.05	82.33 $\pm$ 4.64	81.33 $\pm$ 1.25	68.00 $\pm$ 1.41 [‡]
0.008	98.67 $\pm$ 0.94 [‡]	85.33 $\pm$ 2.62	83.00 $\pm$ 4.32	60.33 $\pm$ 0.94

Notation. [†] marks the clean point-best in that task column; [‡] marks the px+goal 0.08 point-best. Table 4 and Figure 4 use a separate term, *representative diagnostic checkpoint*, for the checkpoint on which the full diagnostic suite was executed. Because many neighbouring settings differ by ≤ 3 pts and overlap in across-seed std, we treat the optima as plateaus unless the gap is clearly larger than the seed-level variability.

**Figure 2:** Noise-training sweep: clean vs OOD per task. Dashed vertical lines mark each task’s px+g 0.08 point-best σ^* . No single **std_max** is jointly optimal across tasks.

Three observations follow.

(1) No single `std_max` is jointly optimal; clean and robust optima can dissociate within a task. TwoRoom peaks globally at `std_max` = 0.008 (98.33 / 98.67); clean rises monotonically with noise. PushT peaks on clean at `std_max` = 0.003 but on robustness at `std_max` = 0.006 (87.00 vs. 0.002’s 71.33; +15.67 pt). Reacher lies on a 0.002–0.006 plateau: clean point-best is at `std_max` = 0.006 (86.00), while px+g 0.08 point-best is at `std_max` = 0.002 (85.67). At `std_max` = 0.001 clean is 61.67 vs base 58.67 — within across-seed std (~ 2.5 pts), so the data support “*low noise is statistically equivalent to base*”, not “low noise hurts”. Cube responds least: clean is non-monotonic, and px+g 0.08 improves only in the 0.003–0.007 range.

(2) Per-task tuning is necessary, not optional. Optimal `std_max` varies substantially: TwoRoom clean/OOD point-best 0.008, PushT clean point-best 0.003 / px+g 0.08 point-best 0.006, Reacher clean point-best 0.006 / px+g 0.08 point-best 0.002, Cube px+g 0.08 point-best 0.007 with a shallow clean plateau around 0.001 / 0.004 / 0.007. Global input-side noise is the strongest “global” form of invariance pressure, but closing the OOD gap requires per-task tuning cost.

(3) Tasks form a clear sensitivity ordering at the extremes. PushT is clearly most sensitive (−81.67 base drop), Cube least sensitive (−20.33), while TwoRoom (−44.00) and Reacher (−43.67) are effectively tied around a 44-pt drop. Recovery does not scale with sensitivity: TwoRoom recovers most fully (+48.67 pt at `std_max` = 0.008), Cube least (+21.67 pt). Input-side global noise is most effective on visually redundant tasks and gives diminishing returns on structured manipulation.

Plotting the same sweep as a (clean, OOD) trajectory per task (Figure 3) makes the trade-off visually explicit: each task’s sweep curve moves from base far below the $y = x$ diagonal toward the upper-right, with per-task curvature determined by how much OOD gain a task can buy from a given drop in clean.




Figure 3: Per-task Pareto trajectory of (clean, OOD) under noise sweep. Open marker = base; solid markers = `std_max` sweep (colour-by-`std_max`); ringed marker = `px+g 0.08` point-best. Dashed line $y = x$.

4.4 Diagnostic analysis: why global noise is not a silver bullet

Table 4 compares core diagnostic metrics on LeWM-base versus the per-task representative noise-trained diagnostic checkpoint.

Table 4: Representation diagnostics: LeWM-base vs. a representative noise-trained diagnostic checkpoint per task. The σ pick per task (0.008 / 0.002 / 0.006 / 0.001) is the ckpt on which the diagnostic suite was originally executed. Under the unified 3-seed $\times$ 100 eval protocol the PushT clean point-best shifts to `std_max` = 0.003 (within ± 2 pt of `std_max` = 0.002); we keep the diagnostics on the `std_max` = 0.002 checkpoint because the compression-vs.-resolution pattern this table illustrates is robust within this neighbourhood.

Metric	TwoRoom base	TwoRoom noise (0.008)	PushT base	PushT noise (0.002)	Reacher base	Reacher noise (0.006)	Cube base	Cube noise (0.001)
<code>clean_nn_cos_dist_median</code>	0.0449	0.0281	0.2360	0.1051	0.0633	0.0676	0.1856	0.1879
<code>clean_effective_rank</code>	47.60	33.59	76.42	42.85	61.04	65.92	73.25	71.83
<code>transition_resolution_ratio_l2</code>	0.7216	0.6055	0.3015	0.2800	0.3704	0.3791	0.4847	0.4629
<code>transition_resolution_ratio_cos</code>	0.5538	0.3780	0.0868	0.0800	0.1351	0.1399	0.2347	0.2168
<code>id_probe_r2</code>	0.2889	-0.0573	0.7739	0.7500	0.1621	0.1729	0.6657	0.6720
<code>action_mean_pred_shift_norm</code>	0.5329	0.4482	0.1283	0.1200	0.2518	0.2585	0.2364	0.2320
<code>predictor_rollout_T8_l2</code>	18.62	17.90	18.65	16.50	15.17	0.44	20.20	19.25

Notes. (i) the TwoRoom transition-resolution values (L2 and cos) are corrected to the values directly stored in the released geometry-summary and task-resolution JSONs; (ii) Reacher and Cube representative diagnostics are taken from the corresponding checkpoint’s per-tool JSON

outputs under the diagnostics directory ($\text{max-std} = 0.1$, history-only noise); (iii) the Reacher representative rollout drift of 0.44 (a $35\times$ reduction from base 15.17) is not a recording error: Reacher’s representative checkpoint trains under $\text{std_max} = 0.006$, which is enough to collapse the multi-step predictor drift while keeping single-step task-resolution metrics flat — precisely the dissociation that motivates the analysis in Section 4.5. Cube’s representative drift ($19.25 \approx \text{base } 20.20$) is the opposite extreme.




Figure 4: Per-task diagnostic radar: base vs representative noise-trained diagnostic checkpoint on 6 metrics, min-max normalised across all 4 tasks.

Mechanistic reading.

- **TwoRoom.** Low-dimensional, discrete, visually redundant. Compressing the representation (effective rank $47.6 \rightarrow 33.6$) is acceptable and even beneficial: a more compact latent space is easier to plan in.
- **PushT.** Continuous contact requires fine-grained pose resolution. Even at the representative diagnostic checkpoint ($\text{std_max} = 0.002$), the L2 transition-resolution metric already trends

slightly downward. At heavier noise (e.g. 0.006) it drops further, erasing contact-transition keyframes.

- **Reacher.** Low-dimensional continuous reaching does not show a resolution collapse in Table 4: effective rank, transition resolution, and the controllability probe remain flat or slightly improve. The notable change is the large reduction in multi-step predictor drift ($15.17 \rightarrow 0.44$), matching the later finding that Reacher’s residual OOD-drop signal lives in the rollout metric rather than in the single-step fragility ratio.
- **Cube.** Cube is moderately resolution-demanding but less fragile than PushT. The representative checkpoint leaves rank, transition resolution, controllability probe, and rollout drift almost unchanged, consistent with Cube’s smaller visual-ODD cliff and the moderate residual signal of multi-step drift in Table 6.
- **Predictor-rollout caveat.** A drop in multi-step predictor rollout drift is not unambiguously good news. It can also mean the latent has become more *predictable* without being more *controllable*: predictor stability bought by sacrificing resolution.

4.5 Cross-checkpoint correlation analysis

Table 5 reports task-specific cross-checkpoint correlations on the LeWM $n = 9$ sweep using the unified 3-seed $\times 100$ evaluation.

We analyse two single-value-per-checkpoint diagnostic metrics that have full coverage on all 9 checkpoints per task: the fragility ratio, defined in Section 3.3 as the single-step predictor target shift normalised by the nearest-neighbour cosine distance at the diagnostic’s max-std injection, and the corresponding multi-step rollout L2 drift. Both are released in `assets/paper1_data/canonical_diagnostics_20260517.json`, derived from each checkpoint’s predictor-sensitivity diagnostic JSON, and are pure checkpoint-level quantities (independent of the evaluation protocol).

Table 5: LeWM $n = 9$ sweep: per-task Pearson r / Spearman ρ vs OOD drop (clean – px+g 0.08). Eval values come from the unified 3-seed $\times 100$ protocol.

Metric	TwoRoom (r/ρ)	PushT (r/ρ)	Reacher (r/ρ)	Cube (r/ρ)
<code>predictor_target_to_nn_cos_ratio_at_max_std</code>	+0.96/+0.72	−0.54/−0.77	+0.97/+0.35	+0.36/+0.21
<code>predictor_rollout_T8_l2_at_max_std</code>	+0.89/+1.00	+0.99/+0.88	+0.99/+0.87	+0.92/+0.54

Reading. Unconditional correlations are large because both diagnostic metrics and the OOD drop are driven by the same upstream variable — `std_max`. The relevant test is partial correlation conditioned on `std_max` (Table 6).

Table 6: Partial Spearman $\rho(\text{metric, OOD drop} \mid \text{std_max})$, $n = 9$ per task.

Metric	TwoRoom	PushT	Reacher	Cube
<code>predictor_target_to_nn_cos_ratio_at_max_std</code>	— (rank ties)	+0.06	−0.12	+0.14
<code>predictor_rollout_T8_l2_at_max_std</code>	— (rank ties)	−0.03	+0.79	+0.50

The TwoRoom partial estimates are undefined because saturating success rates produce rank ties that make residualisation singular at $n = 9$. PushT, Reacher and Cube give clean readings: after removing the monotone sweep trend associated with `std_max`, the **fragility metric carries**

essentially no information about the size of the OOD drop on any task; the **multi-step predictor drift** still carries signal on Reacher (+0.79) and weak signal on Cube (+0.50). The Reacher partial ρ is the only non-trivial residual correlation in the entire matrix.

Table 7: Spearman ρ (unconditional) and partial Spearman ρ conditioned on `std_max`, for the fragility ratio on the PushT $n = 9$ LeWM sweep under the unified 3-seed $\times 100$ eval.

Quantity	Spearman ρ
$\rho(\text{std_max}, \text{metric})$	+0.83
$\rho(\text{std_max}, \text{clean})$	0.00
$\rho(\text{std_max}, \text{px}+\text{goal } 0.08)$	+0.82
$\rho(\text{std_max}, \text{OOD drop})$	−0.93
$\rho(\text{metric}, \text{clean})$ unconditional	−0.33
$\rho(\text{metric}, \text{clean}) \mid \text{std_max}$ partial	− 0.59
$\rho(\text{metric}, \text{px}+\text{g } 0.08)$ unconditional	+0.55
$\rho(\text{metric}, \text{px}+\text{g } 0.08) \mid \text{std_max}$ partial	− 0.41
$\rho(\text{metric}, \text{OOD drop})$ unconditional	−0.77
$\rho(\text{metric}, \text{OOD drop}) \mid \text{std_max}$ partial	+0.06

PushT $n = 9$ detailed view (fragility ratio). Interpretation:

1. **After removing the monotone trend associated with `std_max`, the metric still carries a meaningful residual checkpoint-quality signal.** Partial ρ on clean is −0.59 and on `px+goal 0.08` is −0.41 — both negative, both meaningful at this sample size. On the $n = 9$ PushT sweep, lower fragility ratio aligns with better clean *and* better noisy success beyond what the sweep-level `std_max` trend already explains.
2. **The metric does NOT predict the clean–OOD gap.** The unconditional $\rho(\text{metric}, \text{drop}) = -0.77$ looks impressive, but partialling out `std_max` collapses it to +0.06 (sign-flipped, near zero). The drop’s strong correlation with the metric is a mediated effect: `std_max` drives both the metric ($\rho = +0.83$) and the drop ($\rho = -0.93$). After the `std_max` trend is removed, the metric tells you nothing new about how much the *gap* between clean and noisy performance will be.
3. **Note the unconditional-vs-partial sign reversal on clean.** $\rho(\text{metric}, \text{clean})$ is only −0.33 unconditionally — clean success rate is roughly flat across the PushT sweep under the 3-seed protocol (range 80.67–89.67), so `std_max` barely moves clean. The partial correlation *strengthens* to −0.59 once the sweep-level `std_max` trend is removed, exactly because the residual signal is what the metric tracks.
4. **Practical reading.** The toolkit is a model-selection tool that ranks checkpoints after controlling for the sweep-level `std_max` effect. It is *not* a substitute for actually training with noise when the OOD gap is the quantity of interest.

4.5.1 Clean control fidelity and OOD robustness as separable signals

The partial-correlation analysis above settles the interpretation:

- The metric is a **residual checkpoint-quality signal beyond the sweep-level `std_max` effect** — a lower fragility ratio means better control on *both* clean and noisy evaluation (partial $\rho = -0.59$ / −0.41 on PushT).
- The metric **does not isolate noise robustness** — its apparent correlation with the gap between clean and noisy success is fully mediated by `std_max` (partial $\rho = +0.06$).

The PushT scatter (Figure 5) makes both halves visible. Panel (a) plots metric vs clean (unconditional Spearman $\rho = -0.33$; the residual trend after conditioning on `std_max` is what the partial correlation captures). Panel (b) plots metric vs OOD drop (unconditional $\rho = -0.77$), but the colour-bar (encoding `std_max`) reveals the structure: low-`std_max` checkpoints (light blue) live at high drop, high-`std_max` checkpoints (dark blue) live at low drop, and the metric tracks `std_max` itself. After the `std_max` trend is removed, the metric does not separate small-drop from large-drop checkpoints.




Figure 5: PushT $n = 9$ LeWM sweep: the fragility ratio is a checkpoint-quality predictor (a); the apparent OOD-drop correlation in (b) is mediated by `std_max` (encoded by the colour bar).

4.6 Mechanism attribution: where in the pipeline does noise cause failure?

Section 4.4 reports *what* is compressed and Section 4.5 reports *which diagnostics* correlate with eval across checkpoints, but neither answers “**is the failure in the encoder, the predictor, or the cost surface?**” We address this with two complementary experiments.

4.6.1 Auxiliary cost-swap sanity check: cost alone is unlikely to explain the collapse

We ran a one-off eval-only cost swap on a single TwoRoom checkpoint outside the 36-checkpoint canonical evaluation table; full details are reported in Section E. Switching the CEM cost from cosine/normalized to mse/raw changes `px+goal 0.03` success only from 36.0 to 42.0, still far below a separate clean reference of 69.7. As a sanity check, this suggests that the cost function alone is unlikely to explain the OOD collapse.

4.6.2 Latent-noise probing: encoder is the principal bottleneck

Directly injecting noise into the latent z (skipping the encoder) decouples encoder contributions from predictor + cost. Comparing diagnostic signals across two injection points (pixels vs. latent z , both history-only noise at max std):

- **PushT.** The multi-step input-space metric `predictor_rollout_T8_12_at_max_std` has unconditional $\rho = +0.88$ vs OOD drop (Table 5), driven primarily by `std_max`; after removing the

monotone `std_max` trend, the partial ρ collapses to -0.03 . The single-step fragility ratio (unconditional $\rho = -0.77$; partial $+0.06$ after removing the same trend) tells the same story. Both encoder–predictor signals are mediated by training noise; once that sweep-level effect is accounted for, neither isolates the OOD drop on PushT.

- **Reacher.** `predictor_rollout_T8_12_at_max_std` has unconditional $\rho = +0.87$ and partial $\rho = +0.79$ vs OOD drop — the only non-trivial residual correlation in the matrix. On Reacher, multi-step predictor drift carries genuine OOD-drop information beyond `std_max`.
- **Cube.** `predictor_rollout_T8_12_at_max_std` has unconditional $\rho = +0.54$ and partial $\rho = +0.50$ — a moderate residual signal. Encoder–predictor drift partially explains Cube’s small but non-zero OOD sensitivity.
- **TwoRoom.** Partial correlations are undefined because clean / drop are saturated across the sweep (rank ties). Unconditional ρ still indicates strong encoder–predictor sensitivity to `std_max`. Table 8 summarises the three-layer attribution.

Table 8: Three-layer attribution by task (LeWM $n = 9$ sweep, unified 3-seed $\times$ 100 eval).

Task	Primary cause	Residual after removing the <code>std_max</code> trend
TwoRoom	encoder dominant (rank-saturated)	n/a
PushT	encoder + single-step predictor (both mediated by <code>std_max</code>)	no residual metric isolates OOD drop
Reacher	encoder + multi-step rollout	multi-step drift residual signal (partial $\rho = +0.79$)
Cube	encoder	moderate residual on multi-step drift (partial $\rho = +0.50$)

The common primary cause across the four tasks is **encoder shift transduced by the predictor**. The auxiliary cost-swap sanity check in Section 4.6.1 suggests that the cost function alone is unlikely to explain the collapse, but we do not treat that single-checkpoint ablation as a task-wide quantitative attribution. This is also the root reason Layer-5 task-resolution metrics carry strong signal in Section 4.4: once the encoder’s latent-neighbourhood structure is corrupted beyond the NN distance scale, downstream predictor and planner are already operating on the wrong neighbourhood.




Figure 6: Mechanism schematic. Pixels perturb the encoder, the predictor transduces that shift into planning-relevant latent drift, and CEM acts on the resulting latent geometry. Quantitative attribution in §4.6.2 uses the two full-coverage LeWM $n = 9$ predictor metrics; after conditioning on `std_max`, Reacher’s multi-step drift is the only non-trivial residual signal.

5 Discussion

5.1 Rethinking the “JEPA invariance” narrative

The data challenge an implicit community assumption: latent prediction alone is not sufficient to confer invariance to visual noise. The invariance a JEPA encoder acquires is *in-distribution invariance* — invariance to visual patterns present in the training distribution. When test-time inputs carry high-frequency pixel noise not seen during training, the topology of the latent space breaks down, the predictor outputs incorrect future states, and the planner fails.

This stands in contrast to VJEPA [15], which reports $R^2 > 0.84$ under a Noisy-TV distractor — but that result is on 1D synthetic signals evaluated by linear probe. **Control-task evaluation (success rate) is a strictly stricter standard than linear-probe R^2 :** linear probes only need enough information for a linear classifier to extract, whereas control tasks demand a representation that supports precise latent-space planning and action optimisation.

5.2 The invariance–resolution trade-off — manifestation and scope

The trade-off has the following four-task signature in this paper:

- **TwoRoom.** Background wall colour / texture is pure visual redundancy; discarding it does not impair navigation.
- **PushT.** The pixel changes during T-block / arm contact carry critical force and pose information; discarding them blinds the planner.
- **Reacher.** Low-dimensional continuous control; visual encoding of joint angle requires moderate resolution.
- **Cube.** Object pose and grasp-point spatial relations need moderate resolution, but Cube itself is largely insensitive to global noise (action sequence is structured; visual–action coupling is predictable).

Task-specificity explains why there is no single optimal noise level.

Scope. Whether the trade-off generalises to other latent world-model architectures is an open question we do not address. Reconstruction-based world models such as DreamerV3 [12] explicitly model observations, while decoder-free latent MPC systems such as TD-MPC2 [13] may exhibit different compression dynamics. ViGMO [23] observes related task-specific noise sensitivity on DMC. EMA-target JEPA (I-JEPA / V-JEPA lineage) and variational JEPA may exhibit different dynamics. This paper’s scope is LeWM + CEM; universalising the trade-off exceeds the present evidence.

5.3 When the diagnostic toolkit works — and when it does not

Three scope conditions should be stated up-front so practitioners know when to use the toolkit and when to fall back to direct evaluation.

Scope 1: the toolkit ranks residual checkpoint quality, not OOD-specific robustness.

Section 4.5.1 shows that on PushT, after conditioning on `std_max`, the strongest cross-checkpoint diagnostic (the fragility ratio) has partial Spearman $\rho = -0.59$ with clean success and $\rho = -0.41$ with px+goal 0.08 success. It does *not* isolate OOD-specific robustness: the partial correlation with the clean-to-OOD drop is only +0.06. Use the toolkit to pick the strongest checkpoint after removing the sweep-level `std_max` trend; do not use it as a substitute for an actual OOD eval if your downstream question is robustness.

Scope 2: tasks with weak per-state controllability variance fall outside the toolkit’s strict regime.

Reacher and TwoRoom do not produce a diagnostic-vs-eval Spearman ρ that survives the partial-correlation criterion. Both tasks lack enough residual spread after accounting for `std_max` to distinguish “good” from “bad” checkpoints via a label-free metric. The toolkit can describe what these models have compressed (Table 4) but cannot predict relative checkpoint quality.

Scope 3: cross-checkpoint diagnostics cannot recover what training did not provide.

The largest determinant of OOD drop is whether the model saw noise during training — not which fragility ratio it happens to score on. This is the structural reason $\rho(\text{metric}, \text{drop})$ is weak even when $\rho(\text{metric}, \text{clean})$ is strong: noise vs no-noise training puts each checkpoint on a completely different (clean, OOD) curve (cf. Figure 3), and no static cross-checkpoint diagnostic can substitute for that training-time choice.

The toolkit therefore has a precise scope: it is a *clean-evaluation auxiliary* that lets you select among checkpoints on tasks with strong per-state controllability variance (PushT, Cube), assuming the training protocol is already fixed. It is not an OOD prediction oracle.

5.4 Practical guidance: how to choose `std_max` on a new task

Our sweep suggests a simple operational recipe:

1. Inspect the clean baseline’s fragility ratio as a screening diagnostic, not as an OOD oracle. Very small values (PushT / Cube regime) indicate that local encoder–predictor shift is small relative to the clean nearest-neighbour scale, but they do not by themselves rank OOD sensitivity.
2. Check `clean_effective_rank`, `transition_resolution_ratio`, and task semantics together. PushT combines high rank and high controllability (`id_probe_r`² = 0.7739), so noise should be swept with clean performance as a guardrail. TwoRoom is visually redundant and retains high transition

separability (`transition_resolution_ratio_l2` = 0.7216), so a wider sweep up to 0.008+ is reasonable.

3. Use *two endpoints* in evaluation (clean and max-noise). Checking only one misses one of the optima: PushT’s clean optimum at 0.003 vs robustness optimum at 0.006 is the clearest example.
4. Under compute budget, start with a coarse 4-level sweep ($\{0.001, 0.003, 0.005, 0.007\}$), then refine around the best clean/OOD candidates. The coarse grid is a screening step, not a guarantee of locating the exact point-best `std_max`.

5.5 Limitations and future directions

Limitation 1 — Two backbone families validated; broader JEPA variants remain open.

The toolkit and the trade-off concept are introduced on LeWM; the PLDM sweep in Section F replicates the task-level signatures and the partial-correlation null on PushT, but the full diagnostic apparatus (effective rank trajectory, transition resolution, etc.) is reported only on LeWM in the main text. EMA-target JEPA variants (I-JEPA / V-JEPA lineage) and variational JEPA may exhibit different noise responses; we did not run them.

Limitation 2 — Gaussian pixel noise only. Real-world visual corruption includes motion blur, contrast variation, occlusion, and lighting change; whether the trade-off transfers is open.

Limitation 3 — Diagnostic framework is empirical, not theoretical. Reacher and TwoRoom fail the partial-correlation criterion for *all* metrics, exposing where the empirical framework breaks down.

Limitation 4 — Automated transition reweighting is not a substitute for noise training.

As a sanity check we also tested a scale-preserving heteroscedastic negative-log-likelihood (NLL) formulation in which a learned per-transition σ -head down-weights high-error transitions. The result is informative as a negative finding: TwoRoom clean reaches 99.67% but high-noise robustness lags noise training; PushT clean *collapses* from 86.33% to 13.33% because contact-control transitions have high prediction error and are exactly the transitions σ -weighting would discard. Full data are reported in Section D. The lesson — “hard $\neq$ unimportant” in contact control — connects this paper’s trade-off to the broader question of per-token controllers, which we leave to follow-up work.

Future direction 1 — Per-token adaptive consistency. Sections 4.5 and 4.6 identify the strongest signal — the fragility ratio — as a checkpoint-level scalar. Whether its per-token variant can serve as a per-token consistency controller is an open methodological question, under investigation in follow-up work.

Future direction 2 — Cross-architecture replication. Running the sweep and diagnostic protocol on external world-model baselines would reveal whether the trade-off is JEPA-specific or shared by broader latent world-model families.

Future direction 3 — Theoretical grounding. Reformulating the trade-off in an information-bottleneck / rate-distortion language is an attractive direction; we did not pursue it because the empirical phenomenology may not yet be stable enough for formal modelling.

6 Conclusion

This paper provides a systematic diagnostic study of the visual-OOD control robustness of JEPA + CEM world models, using LeWorldModel as the representative system. Three findings:

1. Visual OOD collapse in JEPA + CEM control. Without noise-aware training, LeWM collapses under pixel noise; latent prediction alone does not confer visual robustness.
2. Global noise augmentation has a boundary. It effectively closes the OOD gap but has no globally optimal level; task differences are large, and within a single task clean and robustness optima can dissociate.
3. A five-layer diagnostic protocol reveals the underlying mechanism. Noise-augmentation gains arise from representation compression, but excessive compression destroys the resolution required for control — the *invariance-resolution trade-off*. When used cross-checkpoint, the strongest single diagnostic (the fragility ratio) predicts clean control quality rather than OOD-specific robustness; we delineate this scope explicitly.

This paper does not propose a new training algorithm. Its contribution is a systematic body of empirical evidence, a reusable diagnostic toolkit, and an explicit delineation of what cross-checkpoint diagnostics can and cannot predict.

Acknowledgements

The complete code and data are available at https://github.com/qun-team/wm_exp. We thank the LeWM authors for releasing the baseline framework on which this study builds.

References

- [1] Guillaume Alain and Yoshua Bengio. Understanding intermediate layers using linear classifier probes, 2017.
- [2] Mahmoud Assran, Quentin Duval, Ishan Misra, Piotr Bojanowski, Pascal Vincent, Michael Rabbat, Yann LeCun, and Nicolas Ballas. Self-supervised learning from images with a joint-embedding predictive architecture. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, 2023.
- [3] Mido Assran, Adrien Bardes, David Fan, Quentin Garrido, Russell Howes, Mojtaba Komeili, Matthew Muckley, Ammar Rizvi, Claire Roberts, Koustuv Sinha, Artem Zholus, Sergio Arnaud, Abha Gejji, Ada Martin, Francois Robert Hogan, Daniel Dugas, Piotr Bojanowski, Vasil Khalidov, Patrick Labatut, Francisco Massa, Marc Szafraniec, Kapil Krishnakumar, Yong Li, Xiaodong Ma, Sarath Chandar, Franziska Meier, Yann LeCun, Michael Rabbat, and Nicolas Ballas. V-jepa 2: Self-supervised video models enable understanding, prediction and planning. *arXiv preprint arXiv:2506.09985*, 2025.
- [4] Adrien Bardes, Quentin Garrido, Jean Ponce, Xinlei Chen, Michael Rabbat, Yann LeCun, Mahmoud Assran, and Nicolas Ballas. Revisiting feature prediction for learning visual representations from video (v-jepa). *Transactions on Machine Learning Research / arXiv:2404.08471*, 2024.
- [5] Adrien Bardes, Jean Ponce, and Yann LeCun. Vicreg: Variance-invariance-covariance regularization for self-supervised learning. In *Proceedings of the International Conference on Learning Representations (ICLR)*, 2022.

- [6] Ting Chen, Simon Kornblith, Mohammad Norouzi, and Geoffrey Hinton. A simple framework for contrastive learning of visual representations. In *Proceedings of the International Conference on Machine Learning (ICML)*, 2020.
- [7] Ekin D. Cubuk, Barret Zoph, Jonathon Shlens, and Quoc V. Le. Randaugment: Practical automated data augmentation with a reduced search space. In *Proceedings of NeurIPS*, 2020.
- [8] T. W. Epps and Lawrence B. Pulley. A test for normality based on the empirical characteristic function. *Biometrika*, 70(3):723–726, 1983.
- [9] Carlos E. García, David M. Prett, and Manfred Morari. Model predictive control: Theory and practice — a survey. *Automatica*, 25(3):335–348, 1989.
- [10] Quentin Garrido, Randall Balestriero, Laurent Najman, and Yann LeCun. Rankme: Assessing the downstream performance of pretrained self-supervised representations by their rank. In *Proceedings of the International Conference on Machine Learning (ICML)*, 2023.
- [11] Hafez Ghaemi, Eilif Muller, and Shahab Bakhtiari. seq-jepa: Autoregressive predictive learning of invariant-equivariant world models. *arXiv preprint arXiv:2505.03176*, 2025.
- [12] Danijar Hafner, Jurgis Pasukonis, Jimmy Ba, and Timothy Lillicrap. Mastering diverse control tasks through world models. *Nature*, 640:647–653, 2025.
- [13] Nicklas Hansen, Hao Su, and Xiaolong Wang. Td-mpc2: Scalable, robust world models for continuous control. In *Proceedings of the International Conference on Learning Representations (ICLR)*, 2024.
- [14] Nicklas Hansen and Xiaolong Wang. Generalization in reinforcement learning by soft data augmentation. In *Proceedings of the IEEE International Conference on Robotics and Automation (ICRA)*, 2021.
- [15] Yongchao Huang. Vjepa: Variational joint embedding predictive architectures as probabilistic world models. *arXiv preprint arXiv:2601.14354*, 2026.
- [16] Li Jing, Pascal Vincent, Yann LeCun, and Yuandong Tian. Understanding dimensional collapse in contrastive self-supervised learning. In *Proceedings of the International Conference on Learning Representations (ICLR)*, 2022.
- [17] Simon Kornblith, Mohammad Norouzi, Honglak Lee, and Geoffrey Hinton. Similarity of neural network representations revisited. In *Proceedings of the International Conference on Machine Learning (ICML)*, 2019.
- [18] Ilya Kostrikov, Denis Yarats, and Rob Fergus. Image augmentation is all you need: Regularizing deep reinforcement learning from pixels. In *Proceedings of the International Conference on Learning Representations (ICLR)*, 2021.
- [19] Dirk P. Kroese, Sergey Porotsky, and Reuven Y. Rubinstein. The cross-entropy method for continuous multi-extremal optimization. In *Methodology and Computing in Applied Probability*, volume 8, pages 383–407, 2006.
- [20] Yann LeCun. A path towards autonomous machine intelligence. OpenReview preprint, 2022.

- [21] Lucas Maes, Quentin Le Lidec, Dan Haramati, Nassim Massaudi, Damien Scieur, Yann LeCun, and Randall Balestriero. *stable-worldmodel-v1: Reproducible world modeling research and evaluation*, 2026.
- [22] Lucas Maes, Quentin Le Lidec, Damien Scieur, Yann LeCun, and Randall Balestriero. *Leworld-model: Stable end-to-end joint-embedding predictive architecture from pixels. arXiv preprint arXiv:2603.19312*, 2026.
- [23] Mingyu Park, Samyeul Noh, Hyun Myung, and Donghwan Lee. Zero-shot visual generalization in model-based reinforcement learning via latent consistency. OpenReview (ICLR 2026 submission), 2025.
- [24] Yuping Qiu, Rui Zhu, and Ying-cong Chen. Improving joint embedding predictive architecture with diffusion noise (n-jepa). *arXiv preprint arXiv:2507.15216*, 2025.
- [25] Ashwath Radhachandran, Vedrana Ivezić, Shreeram Athreya, Ronit Anilkumar, Corey W. Arnold, and William Speier. Us-jepa: A joint embedding predictive architecture for medical ultrasound. *arXiv preprint arXiv:2602.19322*, 2026.
- [26] Olivier Roy and Martin Vetterli. The effective rank: A measure of effective dimensionality. In *Proceedings of the European Signal Processing Conference (EUSIPCO)*, 2007.
- [27] Vlad Sobal, Jyothir S V, Siddhartha Jalagam, Nicolas Carion, Kyunghyun Cho, and Yann LeCun. Joint embedding predictive architectures focus on slow features, 2022.
- [28] Vlad Sobal, Wancong Zhang, Kyunghyun Cho, Randall Balestriero, Tim G. J. Rudner, and Yann LeCun. Stress-testing offline reward-free reinforcement learning: A case for planning with latent dynamics models. In *7th Robot Learning Workshop: Towards Robots with Human-Level Abilities*, 2025.
- [29] Yiyu Sun, Yifei Ming, Xiaojin Zhu, and Yixuan Li. Out-of-distribution detection with deep nearest neighbors. In *Proceedings of the International Conference on Machine Learning (ICML)*, pages 20827–20840, 2022.
- [30] Alex Tamkin, Margalit Glasgow, Xiluo He, and Noah Goodman. Feature dropout: Revisiting the role of augmentations in contrastive learning. In *Proceedings of NeurIPS*, 2023.
- [31] Jayden Teoh, Manan Tomar, Kwangjun Ahn, Edward S. Hu, Pratyusha Sharma, Riashat Islam, Alex Lamb, and John Langford. Next-latent prediction transformers learn compact world models. *arXiv preprint arXiv:2511.05963*, 2025.
- [32] Leonardo F. Toso, Davit Shadunts, Yunyang Lu, Nihal Sharma, Donglin Zhan, Nam H. Nguyen, and James Anderson. Learning invariant visual representations for planning with joint-embedding predictive world models. *arXiv preprint arXiv:2602.18639*, 2026.
- [33] Tongzhou Wang and Phillip Isola. Understanding contrastive representation learning through alignment and uniformity on the hypersphere. In *Proceedings of the International Conference on Machine Learning (ICML)*, 2020.
- [34] Denis Yarats, Rob Fergus, Alessandro Lazaric, and Lerrel Pinto. Mastering visual continuous control: Improved data-augmented reinforcement learning. In *Proceedings of the International Conference on Learning Representations (ICLR)*, 2022.

- [35] Junbo Zhang and Kaisheng Ma. Rethinking the augmentation module in contrastive learning: Learning hierarchical augmentation invariance with expanded views. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, pages 16650–16659, 2022.

A Experimental details

A.1 Environments

- **PushT.** 2D continuous pushing; 20,000 expert episodes; ~ 196 steps/episode; action dim 2 (direction + magnitude); 224×224 RGB.
- **TwoRoom.** 2D continuous navigation; 10,000 episodes; ~ 92 steps; action dim 2; 224×224 RGB.
- **Reacher.** 2D arm control (DeepMind Control Suite); 10,000 episodes; 200 steps; action dim 2.
- **Cube.** 3D cube manipulation (OGBench); 10,000 episodes; 200 steps; action dim 7.

A.2 Noise augmentation implementation

The actual implementation is `utils.py::AddNormalizedGaussianNoise`. Key points: (i) noise is added on ImageNet-normalized tensors and must be divided by the per-channel std to retain “pixel-space std” semantics; (ii) sampling is per-frame independent over the leading frame dims, not per-batch; (iii) all sweeps in this paper fix `noise_prob = 1.0` and `std_min = 0` and vary only `std_max`.

```
class AddNormalizedGaussianNoise:
    """Per-frame independent. Each frame draws Bernoulli(noise_prob) then
    std ~ Uniform(std_min, std_max). Pixel-space std is converted to
    normalized space by dividing by per-channel std before adding."""
    def __init__(self, std_min, std_max, noise_prob=1.0):
        self.std_low, self.std_high, self.noise_prob = std_min, std_max, noise_prob
        self.channel_std = torch.as_tensor(IMAGENET_STD) # (C,)

    def __call__(self, x):
        if self.std_high <= 0 or self.noise_prob <= 0:
            return x
        leading = x.shape[:-3]
        stds = torch.empty(leading, device=x.device, dtype=x.dtype).uniform_(
            self.std_low, self.std_high)
        if self.noise_prob < 1.0:
            mask = (torch.rand(leading, device=x.device) < self.noise_prob).to(x.dtype)
            stds = stds * mask
        per_frame_scale = stds.view(*leading, 1, 1, 1)
        channel_factor = (1.0 / self.channel_std.to(x.device, x.dtype)).view(
            *([1] * len(leading)), -1, 1, 1)
        scale = per_frame_scale * channel_factor # pixel-space std -> normalized
        return x + torch.randn_like(x) * scale
```

A.3 Evaluation protocol

Clean evaluation uses no noise, 100 trajectories per seed $\times 3$ seeds (42/43/44), totalling 300 trajectories per condition. Noisy evaluation adds Gaussian noise of `std` $\in \{0.05, 0.08\}$ to both pixels and the goal image under the same 3-seed protocol. Success criterion is task-specific (PushT: T-block pose match; TwoRoom: target region; Reacher: joint-angle match; Cube: cube position match).

A.4 Compute

Training runs on a single NVIDIA H800 (80 GB) and follows the LeWM training schedule released with the baseline.

A.5 Main-figure rendering

The six main figures are produced by `tools/paper1_figs.py`; run `python-mtools.paper1_figs--out-dir=assets/paper1_figs`. The script reads `assets/paper1_data/canonical_evals_20260517.json` together with `assets/paper1_data/canonical_diagnostics_20260517.json`; no checkpoint or model loading is required.

B Full diagnostic profile of LeWM-base on four tasks

Table 9 summarises every core diagnostic on the four LeWM-base checkpoints. Data are taken from each checkpoint’s per-tool JSON outputs under `eval_results/diagnostics/` (the geometry, task, predictor and latent_noise families).

Table 9: Full LeWM-base diagnostics across four tasks.

Layer / Metric	TwoRoom	PushT	Reacher	Cube	Unit / note
<i>Encoder geometry</i>					
<code>clean_nn_cos_dist_median</code>	0.0449	0.2360	0.0633	0.1856	cosine distance
<code>clean_pair_cos_dist_median</code>	0.9904	1.0228	1.0252	1.0193	pair-wise cos dist
<code>clean_effective_rank</code>	47.60	76.42	61.04	73.25	effective rank
<i>Noise sensitivity</i>					
<code>noise_angle_deg_median</code> (@std= 0.005)	5.51°	1.33°	3.22°	1.40°	median angular shift
<code>noise_to_nn_cos_ratio_median</code>	0.1031	0.0011	0.0249	0.0016	noise/NN cos ratio
<code>robust_radius_std</code>	0.0142	0.0537	0.0142	0.0356	critical noise std
<code>first_risk_std</code>	>0.08	>0.08	>0.08	>0.08	first high-risk std
<code>noise_angle_slope_deg_per_std</code>	1085.8	284.8	831.7	327.0	°/std angular gain
<code>geometry_flag</code>	balanced	robust	balanced	robust	geometry label
<i>Task resolution</i>					
<code>transition_resolution_ratio_l2</code>	0.7216	0.3015	0.3704	0.4847	L2 resolution ratio
<code>transition_resolution_ratio_cos</code>	0.5538	0.0868	0.1351	0.2347	cos resolution ratio
<code>id_probe_r2</code>	0.2889	0.7739	0.1621	0.6657	linear probe R^2
<code>id_probe_r2_min</code>	0.2599	0.6786	0.1366	0.0972	min probe R^2
<code>lidar_rank</code>	46.06	13.95	45.90	42.46	LiDAR rank proxy
<i>Action effect</i>					
<code>action_mean_pred_shift_norm</code>	0.5329	0.1283	0.2518	0.2364	action-perturb mean shift
<code>action_perturb_pred_shift_corr</code>	0.2847	0.2873	0.4042	0.2559	shift $\times \ a\ $ corr
<i>Predictor stability</i>					
<code>predictor_rollout_T8_l2</code>	18.62	18.65	15.17	20.20	T=8 rollout L2 drift
<code>predictor_target_to_nn_cos_ratio_at_max_std</code>	1.51e-4	3.54e-6	2.67e-5	3.39e-6	max std target/NN ratio
<i>Latent noise</i>					
<code>cka_linear_at_max_std</code>	0.1986	0.5536	0.3085	0.1814	CKA clean vs noisy
<code>latent_cost_surface_slope_z</code>	635.31	1.3886	599.45	0.6208	goal latent cost slope

C Heteroscedastic-loss formulation

The scale-preserving heteroscedastic NLL referenced in Section 5.5 (Limitation 4) and detailed in Section D is:

$$\mathcal{L}_{\text{hetero}} = \frac{1}{2} \exp(-s_t) \|z_{t+1} - \hat{z}_{t+1}\|^2 + \frac{1}{2} s_t, \quad (4)$$

where s_t is the σ -head-predicted log-variance. During training $\exp(-s_t)$ acts as an automatic weight: high-error transitions are down-weighted and low-error transitions are up-weighted. When $s_t \equiv 0$ the loss reduces to plain MSE.

D Heteroscedastic-loss negative result (data)

This appendix records the full data behind Section 5.5 Limitation 4. The σ -head is trained jointly with the predictor; gradient is detached on the σ path so the mean prediction path is exactly LeWM MSE when σ is constant.

Table 10: Heteroscedastic-loss evaluation.

Task / model	Clean	goal 0.05	pixels 0.05	px+goal 0.05	goal 0.08	px+goal 0.08
TwoRoom LeWM-base	94.00	73.33	72.33	61.33	58.67	50.00
TwoRoom LeWM+noise point-best	98.33	98.00	98.33	98.00	98.67	98.67
TwoRoom hetero	99.67	85.33	96.67	84.67	73.33	55.33
PushT LeWM-base	86.33	38.00	17.00	12.00	11.33	4.67
PushT LeWM+noise point-best	89.33	87.67	88.00	88.33	89.67	87.00
PushT hetero	13.33	7.67	7.67	7.67	9.67	6.00

TwoRoom hetero reaches 99.67% clean — consistent with the prior that low-dimensional discrete tasks benefit from stronger invariance / clustering — but lags noise training on high-noise robustness. **PushT hetero clean is 13.33% — a method-level failure**, not a robustness trade-off.

Table 11: Representation diagnostics of the hetero-loss failure.

Metric	TwoRoom base	TwoRoom hetero	PushT base	PushT hetero
clean_nn_cos_dist_median	0.0449	0.0281	0.2360	0.1051
clean_effective_rank	47.60	33.59	76.42	42.85
transition_resolution_ratio_cos	0.5538	0.3780	0.0868	0.0101
transition_resolution_ratio_l2	0.7216	0.6055	0.3015	0.1023
id_probe_r2	0.2889	−0.0573	0.7739	0.2678
action_mean_pred_shift_norm	0.5329	0.4482	0.1283	0.0841
predictor_rollout_T8_l2	18.62	17.90	18.65	14.01

Hetero-loss compresses the representation in both tasks. For TwoRoom (low-dimensional, discrete, redundant), compression is acceptable. For PushT, `transition_resolution_ratio_l2` collapses from 0.3015 to 0.1023 and `id_probe_r2` from 0.7739 to 0.2678 — task-relevant state information is erased. The drop in `predictor_rollout_T8_l2` is not good news either: the latent has become more *predictable* without being more *controllable*.

Mechanism. The σ -head correctly learns per-transition difficulty (high σ on contact frames in PushT, on doorway crossings in TwoRoom, etc.), but using σ as an automatic weight to down-weight

high-error transitions misclassifies the contact-and-control-critical states of PushT as “unimportant” and erases them. This is a direct demonstration of the broader trade-off lesson: in contact-heavy control, **hard** $\neq$ **unimportant**.

This finding motivates a follow-up direction we are currently exploring: per-token *consistency* (not loss-reweighting) controllers driven by detached difficulty signals, which preserve the mean prediction path’s gradient distribution. That work is independent of this paper’s diagnostic study and will be reported separately.

E One-off cost-swap sanity check

This appendix records the eval-only cost-swap ablation referenced in Section 4.6.1. It is **not** part of the 36-checkpoint canonical evaluation table and uses a separate clean reference (`num_eval = 300`), so we report it only as a sanity check rather than a pooled statistic.

Table 12: One-off eval-only cost swap on a TwoRoom checkpoint, `std = 0.03 pixels+goal`. Reference row reports a separate clean eval of the same checkpoint (`num_eval = 300`).

Variant	Cost type	Cost space	Success rate
A (default)	cosine	normalized	36.0
B (swap)	mse	raw	42.0
Clean reference (same ckpt)	—	—	69.7

Swapping cost recovers only +6 pt ($36 \rightarrow 42$), still far below the clean reference (69.7). The narrow reading is therefore: **a sanity check suggests the cost function alone is unlikely to explain the OOD collapse.**

F Cross-method replication: PLDM noise sweep on four tasks

This appendix records the second method-family on which the noise sweep and the cross-checkpoint correlation analysis replicate. PLDM follows the joint-embedding predictive line from Sobal et al. [27]; the concrete latent-dynamics planning baseline is from Sobal et al. [28]. We use the implementation distributed through the `stable-worldmodel` baseline suite [21]. Training and evaluation follow the LeWM canonical protocol bit-for-bit: per-frame Gaussian pixel noise via `utils.py::AddNormalizedGaussianNoise`, `std_max` $\in \{0.001, \dots, 0.008\}$, three evaluation seeds (42/43/44), 100 trajectories per seed.

Coverage at the time of this writing. The PLDM `std_max` $\in \{0.001, \dots, 0.008\}$ sweep is complete on all four tasks (32 checkpoints). A clean PLDM baseline (`std_max = 0`) is currently available only on PushT (from the earlier sanity-check run); the three remaining baselines are still training and are reported in a forthcoming revision. All PLDM canonical data lives in `assets/paper1_data/canonical_evals_pldm_20260522.json` and `assets/paper1_data/canonical_diagnostics_pldm_20260522.json`; we do *not* merge these into the LeWM canonical files used by Tables 1–7.

OOD cliff on PushT. The clean-trained PLDM checkpoint replicates the same direction of failure as LeWM under pixel noise on PushT (Table 13). PLDM starts from a lower clean success on PushT than LeWM-base, but still loses 65.33 pt under `px+goal 0.08`. The implication — that

the control-time visual-noise cliff is not a LeWM-specific artefact — is the same point we made in Section 4.2.

Table 13: Clean-trained PLDM vs. LeWM-base on PushT. Success rate mean $\pm$ population std, 3 seeds $\times$ 100 trajectories.

Model	clean	px+goal 0.05	px+goal 0.08	clean $\rightarrow$ 0.08 drop
LeWM-base	86.33 ± 2.36	12.00 ± 4.55	4.67 ± 2.05	-81.67
PLDM clean-trained	75.33 ± 3.68	43.67 ± 4.64	10.00 ± 2.16	-65.33

PLDM full sweep — clean evaluation. Table 14 reports the per-task clean success rate at every trained `std_max` level. Table 15 reports the corresponding px+goal 0.08 robustness.

Table 14: PLDM+noise sweep, **clean** success rate (%). $\dagger$ marks the per-task clean point-best. PushT row at `std_max` = 0 is the baseline of Table 13; the three other-task baselines are pending.

<code>std_max</code>	TwoRoom	PushT	Reacher	Cube
0 (base)	—	75.33 ± 3.68	—	—
0.001	96.33 ± 0.94	72.33 ± 4.19	78.00 ± 0.82	51.33 ± 1.25
0.002	97.33 ± 0.47	71.33 ± 3.86	82.33 ± 2.36	53.33 ± 1.70
0.003	96.67 ± 1.70	$76.67 \pm 7.54^\dagger$	$83.00 \pm 3.74^\dagger$	52.67 ± 3.30
0.004	97.67 ± 0.47	68.67 ± 5.25	45.00 ± 5.10	54.00 ± 2.94
0.005	97.67 ± 1.25	74.33 ± 3.40	72.00 ± 1.41	54.00 ± 2.16
0.006	98.00 ± 0.82	70.33 ± 6.18	79.00 ± 0.82	$56.00 \pm 4.97^\dagger$
0.007	98.00 ± 0.82	66.67 ± 8.38	80.67 ± 2.62	53.67 ± 3.68
0.008	$98.33 \pm 0.94^\dagger$	68.00 ± 1.41	80.33 ± 1.25	54.00 ± 1.63

Table 15: PLDM+noise sweep, **px+goal 0.08** success rate (%). ‡ marks the per-task px+goal 0.08 point-best.

std_max	TwoRoom	PushT	Reacher	Cube
0 (base)	—	10.00 ± 2.16	—	—
0.001	89.33 ± 3.30	16.67 ± 4.03	78.00 ± 1.63	53.00 ± 1.63
0.002	93.33 ± 0.94	40.00 ± 0.82	79.33 ± 2.05	52.33 ± 3.40
0.003	94.33 ± 1.70	$72.00 \pm 2.94^{\ddagger}$	79.33 ± 2.36	55.67 ± 4.11
0.004	98.00 ± 0.00	62.33 ± 7.41	24.00 ± 5.10	53.00 ± 3.56
0.005	$98.33 \pm 0.94^{\ddagger}$	69.33 ± 4.92	59.67 ± 3.30	53.00 ± 1.63
0.006	97.33 ± 1.25	69.67 ± 5.25	76.00 ± 1.41	53.33 ± 3.09
0.007	96.67 ± 0.47	67.00 ± 9.09	77.67 ± 3.30	52.00 ± 0.00
0.008	97.00 ± 0.82	66.33 ± 4.78	$80.67 \pm 2.36^{\ddagger}$	$56.33 \pm 1.25^{\ddagger}$

Reading. Three quantitative observations carry over cleanly from LeWM:

1. **TwoRoom (visually redundant) prefers heavy noise on both methods.** PLDM clean rises monotonically from 96.33% at `std_max` = 0.001 to 98.33% at `std_max` = 0.008. The same monotone trend appears in LeWM (Table 2). Across-seed std drops to ≤ 1.7 pt on PLDM TwoRoom, indicating a smooth response surface.
2. **PushT (contact-heavy) exhibits the OOD cliff and large noise-training recovery.** PLDM clean-trained drops by 65.33pt under px+goal 0.08 (Table 13); training with `std_max` = 0.003 recovers px+goal 0.08 success to 72.00% (+62.00pt over the clean-trained baseline). The clean point-best at `std_max` = 0.003 coincides with LeWM’s clean point-best.
3. **Cube (structured 3D manipulation) is the least noise-responsive on PLDM as well.** The PLDM Cube clean range across the sweep is 51.33–56.00% ($\Delta = 4.67$ pt), and the px+goal 0.08 range is 52.00–56.33% ($\Delta = 4.33$ pt). Both are well within the across-seed variability we observed on LeWM Cube.

Two findings are method-specific and we report them as such. (i) PLDM’s PushT clean point-best is lower than LeWM’s ($\approx 76\%$ vs $\approx 90\%$); we do not interpret this as evidence either way about the trade-off itself, since PLDM is trained with a different loss family and a different latent dimension. (ii) PLDM’s PushT clean/robust optima are co-located at `std_max` = 0.003, whereas LeWM’s clean and robust optima dissociate (`std_max` = 0.003 vs 0.006). The dissociation phenomenon documented in Section 4.3 therefore replicates as a *family signature* (both methods have a non-trivial clean–robust gap) but the exact location is method-dependent.

Cross-method partial correlation on PushT ($n = 9$ per method). We repeat the within-method partial correlation of Table 7 on the PLDM PushT $n = 9$ sweep (Table 16). The four headline numbers are: $\rho(\text{std_max}, \text{metric}) = +0.93$; unconditional $\rho(\text{metric}, \text{OOD drop}) = -0.78$; partial $\rho(\text{metric}, \text{OOD drop}) \mid \text{std_max} = -0.14$. The negative result (apparent metric→OOD-drop signal is mediated by `std_max`) replicates almost exactly across the two methods.

Table 16: PLDM PushT, $n = 9$ sweep: Spearman ρ (unconditional) and partial Spearman ρ conditioned on `std_max`, for the fragility ratio. Numbers from `assets/paper1_data/cross_method_corr_pldm_20260522.json`.

Quantity	Spearman ρ
$\rho(\text{std_max}, \text{metric})$	+0.93
$\rho(\text{std_max}, \text{clean})$	−0.32
$\rho(\text{std_max}, \text{px+goal } 0.08)$	+0.86
$\rho(\text{std_max}, \text{OOD drop})$	−0.87
$\rho(\text{metric}, \text{clean})$ unconditional	−0.68
$\rho(\text{metric}, \text{clean}) \mid \text{std_max}$ partial	−0.22
$\rho(\text{metric}, \text{px+goal } 0.08)$ unconditional	+0.62
$\rho(\text{metric}, \text{px+goal } 0.08) \mid \text{std_max}$ partial	+0.30
$\rho(\text{metric}, \text{OOD drop})$ unconditional	−0.78
$\rho(\text{metric}, \text{OOD drop}) \mid \text{std_max}$ partial	−0.14

Joint $n = 18$ (LeWM PushT + PLDM PushT). Partial Spearman $\rho(\text{metric}, \text{OOD drop}) \mid (\text{std_max}, \text{method-dummy}) = +0.11$. After removing both the sweep-level `std_max` trend and the cross-method offset, the fragility ratio retains essentially no residual correlation with the OOD drop on PushT — the same null result reported for LeWM in Table 7 reappears in the larger joint sample.

What this strengthens, and what it does not. The replicated partial-correlation null on PushT raises confidence that *Conclusion C4 generalises beyond LeWM*: cross-checkpoint diagnostics on the predictor’s local sensitivity do not, in either method, isolate OOD robustness after the training-noise effect is removed. The within-PLDM partial $\rho(\text{metric}, \text{clean}) \mid \text{std_max}$ on PushT is only -0.22 (vs -0.59 on LeWM), so the secondary, more positive claim — that the metric carries a residual checkpoint-quality signal on clean inputs — is partially method-dependent on PushT and we keep that hedge in the discussion (Section 5.3, Scope 1). The PLDM diagnostic profile also cautions against equating geometry robustness with control robustness: PLDM PushT’s clean-trained checkpoint has `clean_effective_rank` = 130.13, `geometry_flag` = `robust`, `transition_resolution_ratio_l2` = 0.4710, and `id_probe_r2` = 0.7570, yet its control success under `px+goal 0.08` is 10.00% before noise training.